

```
!ls
```

```
sample_data
```

```
# =====
# CSET419 - Lab 5: Baseline CNN for Image-to-Image Translation
# Production-Level Implementation with Google Drive Integration
# =====
```

```
# CELL 1: Mount Google Drive & Setup Environment
#
```

```
from google.colab import drive
drive.mount('/content/drive')

import os
import sys
import json
import time
import random
import numpy as np
from datetime import datetime
from typing import Dict, List, Tuple, Optional
import warnings
warnings.filterwarnings('ignore')

# Create project directory structure in Drive
PROJECT_NAME = "CSET419_Lab5_EncoderDecoder"
BASE_PATH = f"/content/drive/MyDrive/{PROJECT_NAME}"
CHECKPOINT_DIR = f"{BASE_PATH}/checkpoints"
RESULTS_DIR = f"{BASE_PATH}/results"
LOGS_DIR = f"{BASE_PATH}/logs"
CONFIG_PATH = f"{BASE_PATH}/config.json"

# Create directories
```

```
for dir_path in [BASE_PATH, CHECKPOINT_DIR, RESULTS_DIR, LOGS_DIR]:
    os.makedirs(dir_path, exist_ok=True)

print(f"✓ Project mounted at: {BASE_PATH}")
print(f"✓ Checkpoints: {CHECKPOINT_DIR}")
print(f"✓ Results: {RESULTS_DIR}")
print(f"✓ Logs: {LOGS_DIR}")

# Set random seeds for reproducibility
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
```

```
Mounted at /content/drive
✓ Project mounted at: /content/drive/MyDrive/CSET419_Lab5_EncoderDecoder
✓ Checkpoints: /content/drive/MyDrive/CSET419_Lab5_EncoderDecoder/checkpoints
✓ Results: /content/drive/MyDrive/CSET419_Lab5_EncoderDecoder/results
✓ Logs: /content/drive/MyDrive/CSET419_Lab5_EncoderDecoder/logs
```

```
#importing libraries

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, Dataset
from torch.utils.tensorboard import SummaryWriter
import torchvision
import torchvision.transforms as transforms
from torchvision.utils import make_grid, save_image

# Verify GPU
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"✓ Using device: {device}")
if torch.cuda.is_available():
    print(f"✓ GPU: {torch.cuda.get_device_name(0)}")
    print(f"✓ Memory: {torch.cuda.get_device_properties(0).total_memory / 1e9:.2f} GB")
```

```
# Set PyTorch seeds
torch.manual_seed(SEED)
if torch.cuda.is_available():
    torch.cuda.manual_seed(SEED)
    torch.cuda.manual_seed_all(SEED)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = True # Enable for faster training
```

```
✓ Using device: cuda
✓ GPU: Tesla T4
✓ Memory: 15.64 GB
```

```
# [CELL 3: Hyperparameter Configuration (PRO LEVEL)]
# [CELL 3: Hyperparameter Configuration (PRO LEVEL)]
```

```
class HyperParams:
    """Professional hyperparameter configuration with validation"""

    def __init__(self):
        # Data parameters
        self.dataset = 'CIFAR10'
        self.batch_size = 128
        self.num_workers = 4
        self.pin_memory = True

        # Image parameters
        self.image_size = 32
        self.channels = 3
        self.normalize_range = (-1, 1) # Normalize to [-1, 1] as required

        # Model architecture
        self.encoder_blocks = [64, 128, 256, 512] # Progressive feature extraction
        self.bottleneck_dim = 512
        self.use_batch_norm = True
        self.use_dropout = True
        self.dropout_rate = 0.2
```

```
# Training parameters
self.epochs = 100
self.learning_rate = 2e-4
self.weight_decay = 1e-5
self.scheduler_type = 'cosine' # 'step', 'cosine', 'plateau'
self.warmup_epochs = 5

# Loss configuration
self.loss_type = 'combined' # 'mse', 'l1', 'combined'
self.mse_weight = 0.5
self.l1_weight = 0.5
self.ssim_weight = 0.0 # Optional: add SSIM loss for better quality

# Optimization
self.optimizer = 'adamw' # 'adam', 'adamw', 'sgd'
self.beta1 = 0.9
self.beta2 = 0.999
self.gradient_clip = 1.0

# Checkpointing
self.save_freq = 10 # Save every N epochs
self.keep_best = 3 # Keep top N best checkpoints

# Logging
self.log_freq = 50 # Log every N batches
self.sample_freq = 5 # Generate samples every N epochs

# Augmentation (for robustness)
self.use_augmentation = False # Keep False for pure reconstruction

def to_dict(self) -> Dict:
    return {k: v for k, v in self.__dict__.items()}

def save(self, path: str):
    with open(path, 'w') as f:
        json.dump(self.to_dict(), f, indent=4)
```

```
@classmethod
def load(cls, path: str):
    with open(path, 'r') as f:
        config = json.load(f)
    hparams = cls()
    for k, v in config.items():
        setattr(hparams, k, v)
    return hparams

# Initialize hyperparameters
hparams = HyperParams()
hparams.save(CONFIG_PATH)
print("\u2708 Hyperparameters configured and saved")
print(json.dumps(hparams.to_dict(), indent=2))
```

✓ Hyperparameters configured and saved
{

```
"dataset": "CIFAR10",
"batch_size": 128,
"num_workers": 4,
"pin_memory": true,
"image_size": 32,
"channels": 3,
"normalize_range": [
    -1,
    1
],
"encoder_blocks": [
    64,
    128,
    256,
    512
],
"bottleneck_dim": 512,
"use_batch_norm": true,
"use_dropout": true,
"dropout_rate": 0.2,
"epochs": 100,
"learning_rate": 0.0002,
"weight_decay": 1e-05,
```

```
"scheduler_type": "cosine",
"warmup_epochs": 5,
"loss_type": "combined",
"mse_weight": 0.5,
"l1_weight": 0.5,
"ssim_weight": 0.0,
"optimizer": "adamw",
"beta1": 0.9,
"beta2": 0.999,
"gradient_clip": 1.0,
"save_freq": 10,
"keep_best": 3,
"log_freq": 50,
"sample_freq": 5,
"use_augmentation": false
}
```

```
# CELL 4: Data Loading & Preprocessing (CIFAR10 Paired)
#
class PairedCIFAR10Dataset(Dataset):
    """
    Creates paired images for image-to-image translation.
    For this lab, we use: Input = Slightly corrupted image, Target = Original
    This creates a meaningful translation task while using CIFAR10.
    """

    def __init__(self, root: str, train: bool = True, transform=None, noise_factor: float = 0.1):
        self.cifar10 = torchvision.datasets.CIFAR10(
            root=root,
            train=train,
            download=True,
            transform=None # We'll handle transforms manually
        )
        self.transform = transform
        self.noise_factor = noise_factor
        self.train = train
```

```
def __len__(self):
    return len(self.cifar10)

def __getitem__(self, idx):
    img, label = self.cifar10[idx]

    # Convert to tensor [0, 1]
    if self.transform:
        img_tensor = self.transform(img)
    else:
        img_tensor = transforms.ToTensor()(img)

    # Create input by adding noise (denoising task)
    # This is a valid image-to-image translation task
    noise = torch.randn_like(img_tensor) * self.noise_factor
    input_img = torch.clamp(img_tensor + noise, 0, 1)

    # Normalize both to [-1, 1] as required
    normalize = transforms.Normalize((-0.5, -0.5, -0.5), (0.5, 0.5, 0.5))
    input_img = normalize(input_img)
    target_img = normalize(img_tensor)

    return {
        'input': input_img,
        'target': target_img,
        'label': label
    }

# Transforms
train_transform = transforms.Compose([
    transforms.ToTensor(), # Converts [0,255] to [0,1]
])
test_transform = transforms.Compose([
    transforms.ToTensor(),
])
# Create datasets
```

```
train_dataset = PairedCIFAR10Dataset(  
    root='./data',  
    train=True,  
    transform=train_transform,  
    noise_factor=0.1  
)  
  
test_dataset = PairedCIFAR10Dataset(  
    root='./data',  
    train=False,  
    transform=test_transform,  
    noise_factor=0.1  
)  
  
# DataLoaders with optimized settings  
train_loader = DataLoader(  
    train_dataset,  
    batch_size=hparams.batch_size,  
    shuffle=True,  
    num_workers=hparams.num_workers,  
    pin_memory=hparams.pin_memory,  
    drop_last=True,  
    persistent_workers=True if hparams.num_workers > 0 else False  
)  
  
test_loader = DataLoader(  
    test_dataset,  
    batch_size=hparams.batch_size,  
    shuffle=False,  
    num_workers=hparams.num_workers,  
    pin_memory=hparams.pin_memory,  
    persistent_workers=True if hparams.num_workers > 0 else False  
)  
  
print(f"✓ Train samples: {len(train_dataset)}")  
print(f"✓ Test samples: {len(test_dataset)}")  
print(f"✓ Batches per epoch: {len(train_loader)}")
```

```
# Visualize sample
sample = next(iter(train_loader))
print(f"✓ Input range: [{sample['input'].min():.2f}, {sample['input'].max():.2f}]")
print(f"✓ Target range: [{sample['target'].min():.2f}, {sample['target'].max():.2f}]")
```

```
100%|██████████| 170M/170M [00:04<00:00, 39.0MB/s]
```

- ✓ Train samples: 50000
- ✓ Test samples: 10000
- ✓ Batches per epoch: 390
- ✓ Input range: [-1.00, 1.00]
- ✓ Target range: [-1.00, 1.00]

```
# CELL 5: Encoder–Decoder Architecture (Professional)
#
class ConvBlock(nn.Module):
    """Professional Conv Block with BN, Activation, and Dropout"""

    def __init__(self, in_ch: int, out_ch: int, downsample: bool = True,
                 use_bn: bool = True, use_dropout: bool = False, dropout_rate: float = 0.2):
        super().__init__()

        layers = []

        # Convolution
        if downsample:
            layers.append(nn.Conv2d(in_ch, out_ch, 4, stride=2, padding=1, bias=False))
        else:
            layers.append(nn.ConvTranspose2d(in_ch, out_ch, 4, stride=2, padding=1, bias=False))

        # Batch Normalization
        if use_bn:
            layers.append(nn.BatchNorm2d(out_ch))

        # Activation
        layers.append(nn.LeakyReLU(0.2, inplace=True) if downsample else nn.ReLU(inplace=True))
```

```
# Dropout
if use_dropout and downsample:
    layers.append(nn.Dropout2d(dropout_rate))

self.block = nn.Sequential(*layers)

# Skip connection if dimensions match
self.skip = (in_ch == out_ch) and not downsample

def forward(self, x):
    return self.block(x)

class ResidualBlock(nn.Module):
    """Residual Block for better gradient flow"""

    def __init__(self, channels: int, use_bn: bool = True):
        super().__init__()
        self.conv1 = nn.Conv2d(channels, channels, 3, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(channels) if use_bn else nn.Identity()
        self.conv2 = nn.Conv2d(channels, channels, 3, padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(channels) if use_bn else nn.Identity()
        self.relu = nn.ReLU(inplace=True)

    def forward(self, x):
        residual = x
        out = self.relu(self.bn1(self.conv1(x)))
        out = self.bn2(self.conv2(out))
        out += residual
        return self.relu(out)

class EncoderDecoderCNN(nn.Module):
    """
    Professional Encoder-Decoder CNN for Image-to-Image Translation
    Architecture: Input -> Encoder -> Bottleneck -> Decoder -> Output
    """

    def __init__(self, hparams: HyperParams):
```

```
super().__init__()

self.hparams = hparams
channels = hparams.channels
blocks = hparams.encoder_blocks

# ===== ENCODER =====
self.encoder = nn.ModuleList()
in_ch = channels

for i, out_ch in enumerate(blocks):
    self.encoder.append(
        ConvBlock(
            in_ch, out_ch,
            downsample=True,
            use_bn=hparams.use_batch_norm,
            use_dropout=hparams.use_dropout and i < 2, # Dropout only in early layers
            dropout_rate=hparams.dropout_rate
        )
    )
    in_ch = out_ch

# Bottleneck with residual blocks for better representation
self.bottleneck = nn.Sequential(
    ResidualBlock(blocks[-1], hparams.use_batch_norm),
    ResidualBlock(blocks[-1], hparams.use_batch_norm),
    ResidualBlock(blocks[-1], hparams.use_batch_norm),
)

# ===== DECODER =====
self.decoder = nn.ModuleList()
reversed_blocks = list(reversed(blocks))

for i in range(len(reversed_blocks) - 1):
    in_ch = reversed_blocks[i]
    out_ch = reversed_blocks[i + 1]

    self.decoder.append(
```

```
        ConvBlock(
            in_ch, out_ch,
            downsample=False,
            use_bn=hparams.use_batch_norm,
            use_dropout=False
        )
    )

    # Final output layer (no batch norm, tanh activation for [-1, 1] range)
    self.final = nn.Sequential(
        nn.ConvTranspose2d(reversed_blocks[-1], channels, 4, stride=2, padding=1),
        nn.Tanh() # Output in [-1, 1]
    )

    self._initialize_weights()

def _initialize_weights(self):
    """He initialization for better convergence"""
    for m in self.modules():
        if isinstance(m, (nn.Conv2d, nn.ConvTranspose2d)):
            nn.init.kaiming_normal_(m.weight, mode='fan_out', nonlinearity='relu')
            if m.bias is not None:
                nn.init.constant_(m.bias, 0)
        elif isinstance(m, nn.BatchNorm2d):
            nn.init.constant_(m.weight, 1)
            nn.init.constant_(m.bias, 0)

def forward(self, x):
    # Encoder
    skips = []
    for enc in self.encoder:
        x = enc(x)
        skips.append(x)

    # Bottleneck
    x = self.bottleneck(x)

    # Decoder with skip connections (U-Net style)
```

```
for i, dec in enumerate(self.decoder):
    x = dec(x)
    # Add skip connection from encoder
    if i < len(skips) - 1:
        skip = skips[-(i+2)]
        if x.shape == skip.shape:
            x = x + skip # Residual connection

    # Final output
    x = self.final(x)
return x

def get_param_count(self):
    return sum(p.numel() for p in self.parameters() if p.requires_grad)
```

```
# Initialize model
model = EncoderDecoderCNN(hparams).to(device)
print(f"✓ Model initialized")
print(f"✓ Parameters: {model.get_param_count():,}")
print(model)
```

```
# CELL 6: Loss Functions & Metrics (Professional)
#
```

```
class CombinedLoss(nn.Module):
    """Combined MSE + L1 Loss with optional SSIM"""

    def __init__(self, hparams: HyperParams):
        super().__init__()
        self.hparams = hparams
        self.mse = nn.MSELoss()
        self.l1 = nn.L1Loss()

    def forward(self, pred, target):
        loss = 0
```

```
if self.hparams.loss_type in ['mse', 'combined']:
    loss += self.hparams.mse_weight * self.mse(pred, target)

if self.hparams.loss_type in ['l1', 'combined']:
    loss += self.hparams.l1_weight * self.l1(pred, target)

return loss

class Metrics:
    """Track and compute metrics"""

    @staticmethod
    def psnr(pred, target, max_val=2.0): # max_val=2 because range is [-1, 1]
        """Peak Signal-to-Noise Ratio"""
        mse = torch.mean((pred - target) ** 2)
        if mse == 0:
            return float('inf')
        return 20 * torch.log10(max_val / torch.sqrt(mse))

    @staticmethod
    def ssim(pred, target, window_size=11):
        """Structural Similarity Index (simplified)"""
        # Simplified SSIM calculation
        mu1 = torch.mean(pred)
        mu2 = torch.mean(target)
        sigma1 = torch.std(pred)
        sigma2 = torch.std(target)
        sigma12 = torch.mean((pred - mu1) * (target - mu2))

        c1 = 0.01 ** 2
        c2 = 0.03 ** 2

        ssim_val = ((2 * mu1 * mu2 + c1) * (2 * sigma12 + c2)) / \
                   ((mu1**2 + mu2**2 + c1) * (sigma1**2 + sigma2**2 + c2))
        return ssim_val

# Initialize loss
criterion = CombinedLoss(hparams).to(device)
```

```
print(f"\u2708 Loss function: {hparams.loss_type}")
print(f"\u2708 MSE weight: {hparams.mse_weight}, L1 weight: {hparams.l1_weight}")

        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): LeakyReLU(negative_slope=0.2, inplace=True)
    )
)
(bottleneck): Sequential(
    (0): ResidualBlock(
        (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
```

(0): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)

```
(0): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
(1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(2): ReLU(inplace=True)
)
)
(2): ConvBlock(
(block): Sequential(
(0): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
(1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(2): ReLU(inplace=True)
)
)
)
(final): Sequential(
(0): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
(1): Tanh()
)
)
✓ Loss function: combined
✓ MSE weight: 0.5. L1 weight: 0.5
```

```
# CELL 7: Optimizer & Scheduler (Professional)
```

```
def get_optimizer(model, hparams):
    """Configure optimizer based on hyperparameters"""

    if hparams.optimizer == 'adam':
        optimizer = optim.Adam(
            model.parameters(),
            lr=hparams.learning_rate,
            betas=(hparams.beta1, hparams.beta2),
            weight_decay=hparams.weight_decay
        )
    elif hparams.optimizer == 'adamw':
        optimizer = optim.AdamW(
            model.parameters(),
            lr=hparams.learning_rate,
            betas=(hparams.beta1, hparams.beta2),
```

```
        weight_decay=hparams.weight_decay
    )
    elif hparams.optimizer == 'sgd':
        optimizer = optim.SGD(
            model.parameters(),
            lr=hparams.learning_rate,
            momentum=0.9,
            weight_decay=hparams.weight_decay
        )
    else:
        raise ValueError(f"Unknown optimizer: {hparams.optimizer}")

    return optimizer

def get_scheduler(optimizer, hparams):
    """Configure learning rate scheduler"""

    if hparams.scheduler_type == 'step':
        scheduler = optim.lr_scheduler.StepLR(
            optimizer, step_size=30, gamma=0.5
        )
    elif hparams.scheduler_type == 'cosine':
        scheduler = optim.lr_scheduler.CosineAnnealingWarmRestarts(
            optimizer, T_0=10, T_mult=2, eta_min=1e-6
        )
    elif hparams.scheduler_type == 'plateau':
        scheduler = optim.lr_scheduler.ReduceLROnPlateau(
            optimizer, mode='min', factor=0.5, patience=5
        )
    elif hparams.scheduler_type == 'onecycle':
        scheduler = optim.lr_scheduler.OneCycleLR(
            optimizer,
            max_lr=hparams.learning_rate,
            epochs=hparams.epochs,
            steps_per_epoch=len(train_loader),
            pct_start=0.3
        )
    else:
```

```
scheduler = None

return scheduler

optimizer = get_optimizer(model, hparams)
scheduler = get_scheduler(optimizer, hparams)

print(f"✓ Optimizer: {hparams.optimizer}")
print(f"✓ Initial LR: {hparams.learning_rate}")
print(f"✓ Scheduler: {hparams.scheduler_type}")
print(f"✓ Weight decay: {hparams.weight_decay}")

✓ Optimizer: adamw
✓ Initial LR: 0.0002
✓ Scheduler: cosine
✓ Weight decay: 1e-05
```

```
# CELL 8: Checkpoint Manager (Professional)
#
```

```
class CheckpointManager:
    """Professional checkpoint management with best model tracking"""

    def __init__(self, checkpoint_dir: str, keep_best: int = 3):
        self.checkpoint_dir = checkpoint_dir
        self.keep_best = keep_best
        self.best_losses = [] # List of (loss, path) tuples

    def save(self, model, optimizer, scheduler, epoch: int,
            train_loss: float, val_loss: float, is_best: bool = False):
        """Save checkpoint"""

        checkpoint = {
            'epoch': epoch,
            'model_state_dict': model.state_dict(),
            'optimizer_state_dict': optimizer.state_dict(),
```

```
'scheduler_state_dict': scheduler.state_dict() if scheduler else None,
'train_loss': train_loss,
'val_loss': val_loss,
'hyperparameters': hparams.to_dict()
}

# Regular checkpoint
regular_path = f"{self.checkpoint_dir}/checkpoint_epoch_{epoch:03d}.pt"
torch.save(checkpoint, regular_path)
print(f"\u2708 Saved checkpoint: {regular_path}")

# Best checkpoint
if is_best:
    best_path = f"{self.checkpoint_dir}/best_model.pt"
    torch.save(checkpoint, best_path)
    print(f"\u2708 Saved best model: {best_path}")

    # Track best losses
    self.best_losses.append((val_loss, best_path))
    self.best_losses.sort(key=lambda x: x[0])

    # Remove old best checkpoints if exceeding keep_best
    if len(self.best_losses) > self.keep_best:
        # Keep only the best ones
        self.best_losses = self.best_losses[:self.keep_best]

def load_latest(self, model, optimizer, scheduler):
    """Load the most recent checkpoint"""
    checkpoints = [f for f in os.listdir(self.checkpoint_dir)
                  if f.startswith('checkpoint_epoch_')]

    if not checkpoints:
        return None, 0

    # Sort by epoch number
    checkpoints.sort()
    latest = checkpoints[-1]
    path = f"{self.checkpoint_dir}/{latest}"
```

```
        return self.load(path, model, optimizer, scheduler)

    def load_best(self, model, optimizer, scheduler):
        """Load the best checkpoint"""
        path = f"{self.checkpoint_dir}/best_model.pt"
        if os.path.exists(path):
            return self.load(path, model, optimizer, scheduler)
        return None, 0

    def load(self, path: str, model, optimizer, scheduler):
        """Load specific checkpoint"""
        checkpoint = torch.load(path, map_location=device)

        model.load_state_dict(checkpoint['model_state_dict'])
        optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
        if scheduler and checkpoint['scheduler_state_dict']:
            scheduler.load_state_dict(checkpoint['scheduler_state_dict'])

        epoch = checkpoint['epoch']
        print(f"\u2708 Loaded checkpoint from epoch {epoch}\u2709")
        return checkpoint, epoch

    checkpoint_manager = CheckpointManager(CHECKPOINT_DIR, hparams.keep_best)
    print(f"\u2708 Checkpoint manager initialized")
    print(f"\u2708 Keeping top {hparams.keep_best} best models")
```

- ✓ Checkpoint manager initialized
- ✓ Keeping top 3 best models

```
# [CELL 9: Training Loop (Professional with All Features)]
```

```
class Trainer:
    """Professional training manager"""
```

```
def __init__(self, model, criterion, optimizer, scheduler, hparams):
    self.model = model
    self.criterion = criterion
    self.optimizer = optimizer
    self.scheduler = scheduler
    self.hparams = hparams
    self.writer = SummaryWriter(LOGS_DIR)
    self.global_step = 0
    self.best_val_loss = float('inf')
    self.history = {
        'train_loss': [], 'val_loss': [],
        'train_psnr': [], 'val_psnr': [],
        'learning_rates': []
    }

def train_epoch(self, dataloader, epoch):
    """Train for one epoch"""
    self.model.train()
    total_loss = 0
    total_psnr = 0
    num_batches = 0

    pbar = tqdm(dataloader, desc=f"Epoch {epoch}/{self.hparams.epochs} [Train]")

    for batch_idx, batch in enumerate(pbar):
        inputs = batch['input'].to(device)
        targets = batch['target'].to(device)

        # Forward pass
        outputs = self.model(inputs)
        loss = self.criterion(outputs, targets)

        # Backward pass
        self.optimizer.zero_grad()
        loss.backward()

        # Gradient clipping
        if self.hparams.gradient_clip > 0:
```

```
        torch.nn.utils.clip_grad_norm_(
            self.model.parameters(),
            self.hparams.gradient_clip
        )

        self.optimizer.step()

    # Metrics
    with torch.no_grad():
        psnr_val = Metrics.psnr(outputs, targets).item()

    total_loss += loss.item()
    total_psnr += psnr_val
    num_batches += 1

    # Logging
    if batch_idx % self.hparams.log_freq == 0:
        self.writer.add_scalar('Loss/train_batch', loss.item(), self.global_step)
        self.writer.add_scalar('PSNR/train_batch', psnr_val, self.global_step)

    # Update progress bar
    pbar.set_postfix({
        'loss': f'{loss.item():.4f}',
        'psnr': f'{psnr_val:.2f}',
        'lr': f'{self.optimizer.param_groups[0]['lr']):.6f}'
    })

    self.global_step += 1

    # Update scheduler if OneCycle
    if isinstance(self.scheduler, optim.lr_scheduler.OneCycleLR):
        self.scheduler.step()

    avg_loss = total_loss / num_batches
    avg_psnr = total_psnr / num_batches

return avg_loss, avg_psnr
```

```
@torch.no_grad()
def validate(self, dataloader, epoch):
    """Validation loop"""
    self.model.eval()
    total_loss = 0
    total_psnr = 0
    num_batches = 0

    all_inputs = []
    all_outputs = []
    all_targets = []

    pbar = tqdm(dataloader, desc=f"Epoch {epoch}/{self.hparams.epochs} [Val]")

    for batch in pbar:
        inputs = batch['input'].to(device)
        targets = batch['target'].to(device)

        outputs = self.model(inputs)
        loss = self.criterion(outputs, targets)

        psnr_val = Metrics.psnr(outputs, targets).item()

        total_loss += loss.item()
        total_psnr += psnr_val
        num_batches += 1

        # Store samples for visualization
        if len(all_inputs) < 64: # Store up to 64 samples
            all_inputs.append(inputs.cpu())
            all_outputs.append(outputs.cpu())
            all_targets.append(targets.cpu())

    pbar.set_postfix({
        'loss': f'{loss.item():.4f}',
        'psnr': f'{psnr_val:.2f}'
    })
```

```
# Concatenate samples
all_inputs = torch.cat(all_inputs)[:64]
all_outputs = torch.cat(all_outputs)[:64]
all_targets = torch.cat(all_targets)[:64]

avg_loss = total_loss / num_batches
avg_psnr = total_psnr / num_batches

return avg_loss, avg_psnr, (all_inputs, all_outputs, all_targets)

def visualize_results(self, samples, epoch):
    """Save visualization of results"""
    inputs, outputs, targets = samples

    # Denormalize from [-1, 1] to [0, 1] for visualization
    def denorm(x):
        return (x + 1) / 2

    inputs = denorm(inputs)
    outputs = denorm(outputs)
    targets = denorm(targets)

    # Create grid: [Input | Output | Target]
    comparison = torch.cat([inputs, outputs, targets], dim=0)
    grid = make_grid(comparison, nrow=8, normalize=False, value_range=(0, 1))

    # Save image
    save_path = f"{RESULTS_DIR}/epoch_{epoch:03d}.png"
    save_image(grid, save_path)

    # Add to tensorboard
    self.writer.add_image('Results/Input_Output_Target', grid, epoch)

    return save_path

def train(self, train_loader, val_loader, checkpoint_manager):
    """Full training loop"""
    pass
```

```
print(f"\n{'='*60}")
print(f"Starting Training: {self.hparams.epochs} epochs")
print(f"{'='*60}\n")

start_epoch = 1

# Try to resume from checkpoint
checkpoint, resumed_epoch = checkpoint_manager.load_latest(
    self.model, self.optimizer, self.scheduler
)
if checkpoint:
    start_epoch = resumed_epoch + 1
    self.best_val_loss = checkpoint.get('val_loss', float('inf'))
    print(f"Resuming from epoch {start_epoch}")

for epoch in range(start_epoch, self.hparams.epochs + 1):
    epoch_start_time = time.time()

    # Train
    train_loss, train_psnr = self.train_epoch(train_loader, epoch)

    # Validate
    val_loss, val_psnr, samples = self.validate(val_loader, epoch)

    # Scheduler step (if not OneCycle)
    if self.scheduler and not isinstance(self.scheduler, optim.lr_scheduler.OneCycleLR):
        if isinstance(self.scheduler, optim.lr_scheduler.ReduceLROnPlateau):
            self.scheduler.step(val_loss)
        else:
            self.scheduler.step()

    current_lr = self.optimizer.param_groups[0]['lr']

    # Update history
    self.history['train_loss'].append(train_loss)
    self.history['val_loss'].append(val_loss)
    self.history['train_psnr'].append(train_psnr)
    self.history['val_psnr'].append(val_psnr)
```

```
        self.history['learning_rates'].append(current_lr)

        # TensorBoard logging
        self.writer.add_scalar('Loss/train', train_loss, epoch)
        self.writer.add_scalar('Loss/val', val_loss, epoch)
        self.writer.add_scalar('PSNR/train', train_psnr, epoch)
        self.writer.add_scalar('PSNR/val', val_psnr, epoch)
        self.writer.add_scalar('Learning_Rate', current_lr, epoch)

        # Visualization
        if epoch % self.hparams.sample_freq == 0 or epoch == 1:
            vis_path = self.visualize_results(samples, epoch)

        # Checkpointing
        is_best = val_loss < self.best_val_loss
        if is_best:
            self.best_val_loss = val_loss

        if epoch % self.hparams.save_freq == 0 or is_best or epoch == self.hparams.epochs:
            checkpoint_manager.save(
                self.model, self.optimizer, self.scheduler,
                epoch, train_loss, val_loss, is_best
            )

        # Print epoch summary
        epoch_time = time.time() - epoch_start_time
        print(f"\nEpoch {epoch}/{self.hparams.epochs} Summary:")
        print(f"  Time: {epoch_time:.2f}s | LR: {current_lr:.6f}")
        print(f"  Train Loss: {train_loss:.4f} | PSNR: {train_psnr:.2f}")
        print(f"  Val Loss: {val_loss:.4f} | PSNR: {val_psnr:.2f}")
        print(f"  Best Val Loss: {self.best_val_loss:.4f} {'⭐' if is_best else ''}")
        print("-" * 60)

        # Save final history
        history_path = f"{RESULTS_DIR}/training_history.json"
        with open(history_path, 'w') as f:
            json.dump(self.history, f, indent=4)
```

```
        self.writer.close()
    print(f"\n✓ Training completed! Best Val Loss: {self.best_val_loss:.4f}")
    print(f"✓ Results saved to: {RESULTS_DIR}")
    print(f"✓ Checkpoints saved to: {CHECKPOINT_DIR}")

    return self.history
```

```
# [CELL 10: Initialize & Run Training]
# [REPL 10]

from tqdm import tqdm

# Initialize trainer
trainer = Trainer(model, criterion, optimizer, scheduler, hparams)

# Run training
history = trainer.train(train_loader, test_loader, checkpoint_manager)
```

Epoch 89/100 Summary:

Time: 27.92s | LR: 0.000174
Train Loss: 0.0376 | PSNR: 27.10
Val Loss: 0.0517 | PSNR: 24.95
Best Val Loss: 0.0493

Epoch 90/100 [Train]: 100%|██████████| 390/390 [00:23<00:00, 16.71it/s, loss=0.0376, psnr=27.09, lr=0.000174]
Epoch 90/100 [Val]: 100%|██████████| 79/79 [00:04<00:00, 17.85it/s, loss=0.0570, psnr=24.29]
✓ Saved checkpoint: /content/drive/MyDrive/CSET419_Lab5_EncoderDecoder/checkpoints/checkpoint_epoch_090.pt

Epoch 90/100 Summary:

Time: 28.85s | LR: 0.000171
Train Loss: 0.0374 | PSNR: 27.13
Val Loss: 0.0551 | PSNR: 24.49
Best Val Loss: 0.0493

Epoch 91/100 [Train]: 100%|██████████| 390/390 [00:24<00:00, 16.21it/s, loss=0.0377, psnr=27.03, lr=0.000171]
Epoch 91/100 [Val]: 100%|██████████| 79/79 [00:04<00:00, 17.55it/s, loss=0.0531, psnr=24.77]

Epoch 91/100 Summary:

Time: 28.76s | LR: 0.000168
Train Loss: 0.0375 | PSNR: 27.12
Val Loss: 0.0518 | PSNR: 24.92
Best Val Loss: 0.0493

Epoch 92/100 [Train]: 100%|██████████| 390/390 [00:22<00:00, 17.19it/s, loss=0.0375, psnr=27.06, lr=0.000168]
Epoch 92/100 [Val]: 100%|██████████| 79/79 [00:05<00:00, 14.78it/s, loss=0.0544, psnr=24.64]

Epoch 92/100 Summary:

Time: 28.23s | LR: 0.000165
Train Loss: 0.0373 | PSNR: 27.16
Val Loss: 0.0522 | PSNR: 24.91
Best Val Loss: 0.0493

Epoch 93/100 [Train]: 100%|██████████| 390/390 [00:22<00:00, 17.13it/s, loss=0.0360, psnr=27.43, lr=0.000165]
Epoch 93/100 [Val]: 100%|██████████| 79/79 [00:04<00:00, 16.09it/s, loss=0.0534, psnr=24.76]

Epoch 93/100 Summary:

Time: 27.96s | LR: 0.000162

```
# CELL 11: Evaluation & Analysis  
  
import matplotlib.pyplot as plt  
from PIL import Image  
  
def plot_training_history(history):  
    """Plot training curves"""\n    fig, axes = plt.subplots(2, 2, figsize=(15, 10))  
  
    # Loss curves  
    axes[0, 0].plot(history['train_loss'], label='Train Loss')  
    axes[0, 0].plot(history['val_loss'], label='Val Loss')  
    axes[0, 0].set_title('Loss Curves')  
    axes[0, 0].set_xlabel('Epoch')  
    axes[0, 0].set_ylabel('Loss')  
    axes[0, 0].legend()  
    axes[0, 0].grid(True)  
  
    # PSNR curves  
    axes[0, 1].plot(history['train_psnr'], label='Train PSNR')  
    axes[0, 1].plot(history['val_psnr'], label='Val PSNR')  
    axes[0, 1].set_title('PSNR Curves')  
    axes[0, 1].set_xlabel('Epoch')  
    axes[0, 1].set_ylabel('PSNR (dB)')  
    axes[0, 1].legend()  
    axes[0, 1].grid(True)  
  
    # Learning rate  
    axes[1, 0].plot(history['learning_rates'])  
    axes[1, 0].set_title('Learning Rate Schedule')  
    axes[1, 0].set_xlabel('Epoch')  
    axes[1, 0].set_ylabel('LR')  
    axes[1, 0].set_yscale('log')  
    axes[1, 0].grid(True)  
  
    # Loss difference (overfitting indicator)
```

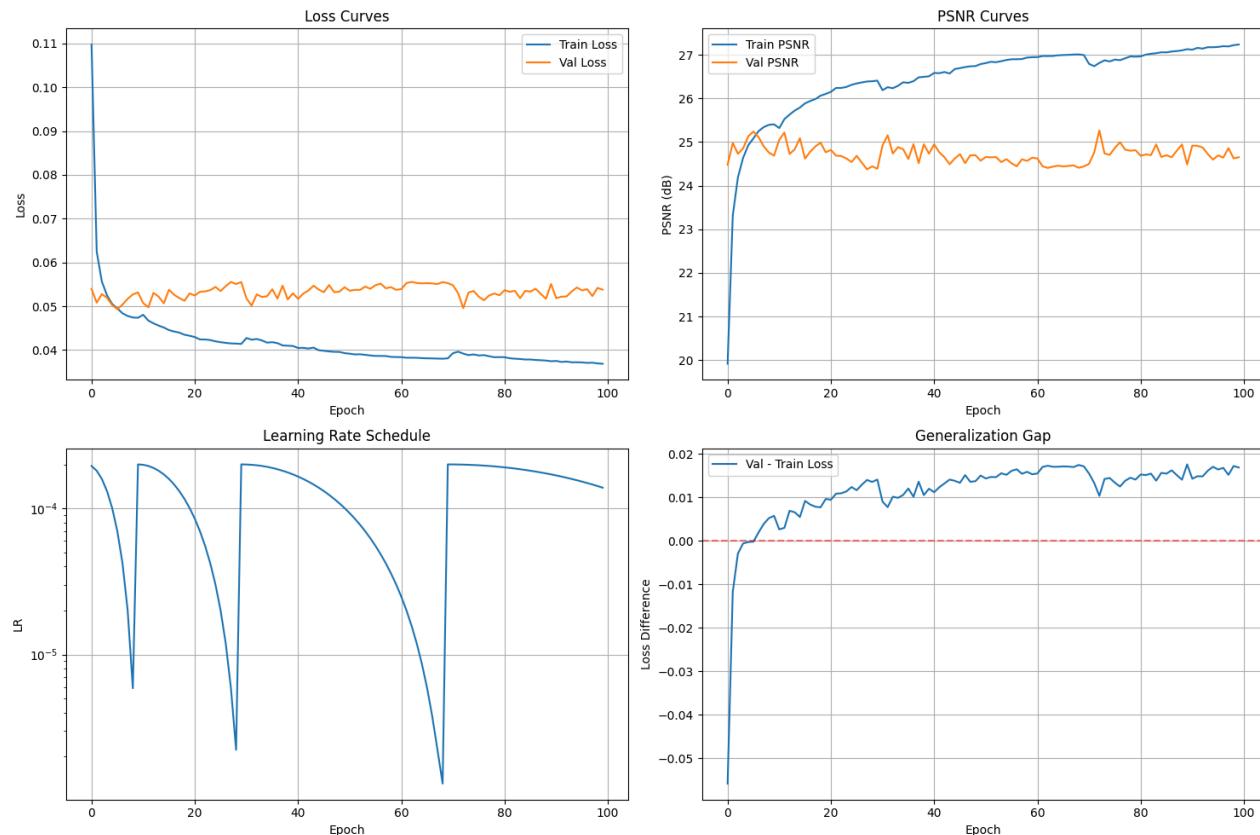
```
diff = np.array(history['val_loss']) - np.array(history['train_loss'])
axes[1, 1].plot(diff, label='Val - Train Loss')
axes[1, 1].axhline(y=0, color='r', linestyle='--', alpha=0.5)
axes[1, 1].set_title('Generalization Gap')
axes[1, 1].set_xlabel('Epoch')
axes[1, 1].set_ylabel('Loss Difference')
axes[1, 1].legend()
axes[1, 1].grid(True)

plt.tight_layout()
plt.savefig(f'{RESULTS_DIR}/training_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

# Plot history
plot_training_history(history)

# Show final results
print("\n" + "="*60)
print("FINAL RESULTS SUMMARY")
print("="*60)
print(f"Best Validation Loss: {min(history['val_loss']):.6f}")
print(f"Best Validation PSNR: {max(history['val_psnr']):.2f} dB")
print(f"Final Training Loss: {history['train_loss'][-1]:.6f}")
print(f"Final Validation Loss: {history['val_loss'][-1]:.6f}")
print("="*60)

# Display sample results
from IPython.display import display
latest_result = f'{RESULTS_DIR}/epoch_{hparams.epochs:03d}.png'
if os.path.exists(latest_result):
    img = Image.open(latest_result)
    plt.figure(figsize=(20, 10))
    plt.imshow(img)
    plt.axis('off')
    plt.title(f'Results at Epoch {hparams.epochs} (Input | Output | Target)')
    plt.show()
```

FINAL RESULTS SUMMARY

Best Validation Loss: 0.049331

Best Validation PSNR: 25.27 dB
Final Training Loss: 0.036855
Final Validation Loss: 0.053773

Results at Epoch 100 (Input | Output | Target)



Start coding or generate with AI.

